

# Food(3) Image Active Learning with Resnet50 on 3 Categories

We are experimenting with 3 categories to improve label performance - Bread, Meat, Veg-Fruits.

Originally, [Food-11](#) is a dataset containing 16643 food images grouped in 11 major food categories. The 11 categories are Bread, Dairy product, Dessert, Egg, Fried food, Meat, Noodles/Pasta, Rice, Seafood, Soup, and Vegetable/Fruit. Similar as Food-5K dataset, the whole dataset is divided in three parts: training, validation and evaluation. The same naming convention is used, where ID 0-10 refers to the 11 food categories respectively.

**Conclusion:** [Uncertainty Sampling](#) out performed all strategies with a **16%** lift in f1 score.

## Experiment results Summary

Overall Classification Metrics Comparison of Base model and Active Learning strategies:

Query Strategy	Initial Model f1 Score	Active Learning f1 Score
Random Sampling	0.71	0.72
Uncertainty Sampling	0.71	0.87
Entropy Sampling	0.71	0.83
Margin Sampling	0.71	0.81

Class-wise Precision Comparison with and without Active Learning:

Query Strategy	Bread	Bread	Meat	Meat	Meat
	Initial Model precision	Active Learning precision			
Random Sampling	0.93	0.95	0.60	0.61	
Uncertainty Sampling	0.93	0.94	0.60	0.79	
Entropy Sampling	0.93	0.94	0.60	0.84	
Margin Sampling	0.93	0.97	0.60	0.70	

Class-wise Recall Comparison with and without Active Learning:

Query Strategy	Bread	Bread	Meat	Meat	Vegetables-Initial Model
	Initial Model recall	Active Learning recall			
Random Sampling	0.34	0.39	0.99	0.99	0.79
Uncertainty Sampling	0.34	0.74	0.99	0.95	0.79
Entropy Sampling	0.34	0.71	0.99	0.87	0.79
Margin Sampling	0.34	0.61	0.99	0.99	0.79

**10% (250 images out of 2500) of training data was used for initial training and  
11% (250 images out of 2250) of remaining for Query Pool.**

Original Training Dataset Size -

1. Bread = 994
2. Meat = 1,325
3. Vegetables-Fruits = 709

Datasize used for this Experiment -

1. Bread = 800
2. Meat = 1,000
3. Vegetables-Fruits = 700

(Bread and Meat images were selected by ranking images by image size and eliminating large sized images)

**This notebook tries to implement -**

Active learning using modAL library - The implementation uses an initial set of labeled examples in order to train an initial model, and then queries unlabeled data to use the best possible examples to train the next model. The notebook uses skorch (scikit-learn compatible neural network library that wraps PyTorch) to handle the training mechanism - to be explored further.

**Querying Strategies:**

1. Uncertainty Sampling: Selects the least sure instances for labelling.
  - a. Classifier uncertainty =  $1 - P(\text{prediction is correct})$
2. Entropy Sampling: Selects the instances where the class probabilities have the largest entropy.
  - a. Entropy of the class probabilities. Example, Confusing predictions with higher probabilities for multiple classes will have higher entropy. Predictions with single dominating probability class will have lower entropy.

3. Margin Sampling: Selects the instances where the difference between the first most likely and second most likely classes are the smallest.

- a. Margin uncertainty, which is the difference of the probabilities of first and second most likely predictions.

Additional resources -

1. <https://dagshub.com/blog/active-learning-your-way-to-better-models/>
2. <https://arxiv.org/pdf/2203.13450.pdf>
3. <https://burrsettles.com/pub/settles.activelearning.pdf>
4. <https://www.youtube.com/watch?v=W2bJH0iXTKc&t=1011s>
5. <https://skorch.readthedocs.io/en/stable/>
6. <https://arxiv.org/pdf/2203.07034.pdf>
7. <https://paperswithcode.com/paper/active-learning-in-annotating-micro-blogs>

## ▼ Active Learning Experiments

### ▼ Load Folder and Libraries

```
from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
!pip install modAL  
!pip install -U skorch
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public  
Collecting modAL  
  Downloading modAL-0.4.1-py3-none-any.whl (27 kB)  
Requirement already satisfied: numpy>=1.13 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: scipy>=0.18 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: pandas>=1.1.0 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from modAL)  
Installing collected packages: modAL  
Successfully installed modAL-0.4.1  
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public
```

```
Collecting skorch
```

```
  Downloading skorch-0.11.0-py3-none-any.whl (155 kB)
```

```
 |██████████| 155 kB 5.0 MB/s
```

```
Requirement already satisfied: tabulate>=0.7.7 in /usr/local/lib/python3.7/dist-packages
```

```
Requirement already satisfied: scikit-learn>=0.19.1 in /usr/local/lib/python3.7/dist-packages
```

```
Requirement already satisfied: tqdm>=4.14.0 in /usr/local/lib/python3.7/dist-packages (1.0.0)
```

```
Requirement already satisfied: scipy>=1.1.0 in /usr/local/lib/python3.7/dist-packages (1.1.0)
```

```
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.7/dist-packages (1.13.3)
```

```
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (1.0.1)
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-packages (2.0.0)
```

```
Installing collected packages: skorch
```

```
Successfully installed skorch-0.11.0
```

```
from modAL.models import ActiveLearner
from modAL.uncertainty import uncertainty_sampling, entropy_sampling, margin_sampling
import numpy as np
import pandas as pd

import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
# import torchvision.transforms as transforms
from torchsummary import summary
import os
import time
import shutil

from imutils import paths
from tqdm import tqdm
import matplotlib.pyplot as plt

import torch
from torch import optim
from torchvision import models
from torchvision import transforms
from torch.cuda.amp import autocast
from torchvision import datasets
from torch.utils.data import DataLoader, Subset
from torchvision.models import densenet121
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from pathlib import Path
import copy

from skorch import NeuralNetClassifier

import warnings
warnings.filterwarnings('ignore')
```

## ▼ Set Parameters

```
# define path to the original dataset
folder = Path('/content/drive/MyDrive/Colab Notebooks/active learning')
DATA_PATH = "/content/drive/MyDrive/Colab Notebooks/active learning/Food-3"

# define base path to store our modified dataset
BASE_PATH = "/content/drive/MyDrive/Colab Notebooks/active learning/Food-3"

# define paths to separate train, validation, and test splits
TRAIN = os.path.join(BASE_PATH, "training")
VAL = os.path.join(BASE_PATH, "validation")
TEST = os.path.join(BASE_PATH, "evaluation")

# initialize the list of class label names
CLASSES = ["Bread", "Meat", "Vegetable-Fruit"]

# specify ImageNet mean and standard deviation and image size
MEAN = [0.485, 0.456, 0.406]
STD = [0.229, 0.224, 0.225]
IMAGE_SIZE = 128

# set the device to be used for training and evaluation
DEVICE = torch.device("cuda")

# specify training hyperparameters
LOCAL_BATCH_SIZE = 128
PRED_BATCH_SIZE = 4
EPOCHS = 20
LR = 0.0001

# determine the number of GPUs we have
NUM_GPU = torch.cuda.device_count()
print(f"[INFO] number of GPUs found: {NUM_GPU}...")

# determine the batch size based on the number of GPUs
BATCH_SIZE = LOCAL_BATCH_SIZE * NUM_GPU
print(f"[INFO] using a batch size of {BATCH_SIZE}...")

# Assign device
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

# Set Seed
torch.manual_seed(42)
np.random.seed(42)

[INFO] number of GPUs found: 1...
[INFO] using a batch size of 128...
```

cuda:0

## ▼ Dataset Transformation and Loaders

```
# define augmentation pipelines
trainTransform = transforms.Compose([
    transforms.RandomResizedCrop(IMAGE_SIZE),
    # transforms.RandomHorizontalFlip(),
    # transforms.RandomRotation(90),
    transforms.ToTensor(),
    transforms.Normalize(mean=MEAN, std=STD)
])

testTransform = transforms.Compose([
    transforms.Resize((IMAGE_SIZE, IMAGE_SIZE)),
    transforms.ToTensor(),
    transforms.Normalize(mean=MEAN, std=STD)
])
```

## ▼ Sampling Original Training data for Experimentation and scoring

```
image_datasets = {
    'train':
        datasets.ImageFolder(TRAIN, transform=trainTransform),
    'validation':
        datasets.ImageFolder(VAL, transform=testTransform),
    'test':
        datasets.ImageFolder(TEST, transform=testTransform),
    'sampler':
        datasets.ImageFolder(TRAIN, transform=trainTransform)
}

image_trainfolder = image_datasets['train']
train_sampler_idx, base_train_idx, target_sampler, base_target = train_test_split(range(len(i

train_ds = Subset(image_trainfolder, base_train_idx)
sampler_ds = Subset(image_trainfolder, train_sampler_idx)

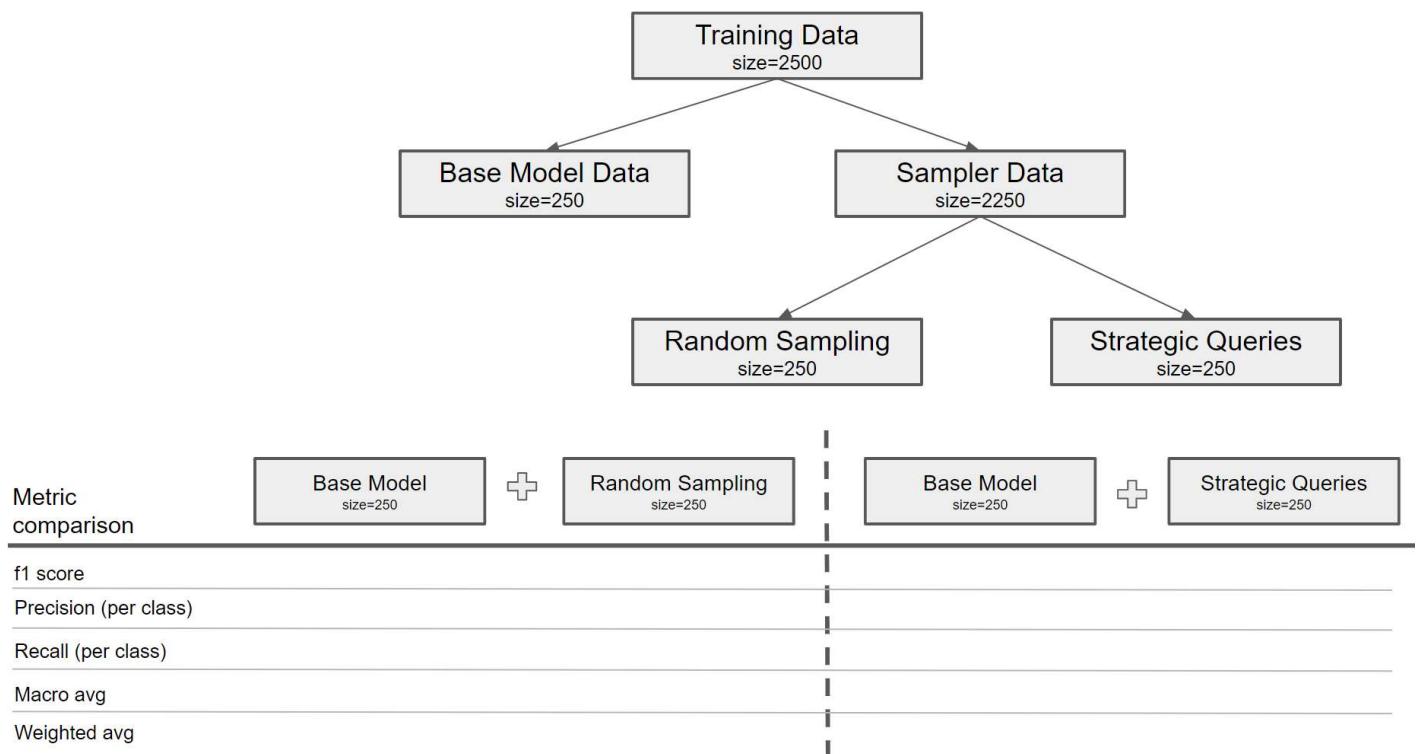
dataloaders = {
    'train':
        torch.utils.data.DataLoader(train_ds #image_datasets['train']
                                    ,batch_size=250,
                                    shuffle=False,
                                    num_workers=0 ),
    'validation':
```

```

    torch.utils.data.DataLoader(image_datasets[ 'validation' ],
                                batch_size=1043,
                                shuffle=False,
                                num_workers=0),
    'test':
        torch.utils.data.DataLoader(image_datasets[ 'test' ],
                                    batch_size=1031,
                                    shuffle=False,
                                    num_workers=0),
    'sampler':
        torch.utils.data.DataLoader(sampler_ds #image_datasets['train']
                                    ,batch_size=2250,
                                    shuffle=False,
                                    num_workers=0)
}

```

## ▼ Experiments Overview



## ▼ RESNET50 Model Training

```

# Load model and define training parameters
model = models.resnet50(pretrained=True).to(device)

for param in model.parameters():

```

```
param.requires_grad = False

model.fc = nn.Sequential(
    nn.Linear(2048, 128),
    nn.ReLU(inplace=True),
    nn.Linear(128, 3)).to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.fc.parameters())

Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/
100% 97.8M/97.8M [00:00<00:00, 223MB/s]
◀ ▶
```

# Define Model training function with 'n' epochs

```
def train_model(model, criterion, optimizer, num_epochs=3):
    for epoch in range(num_epochs):
        print('Epoch {} / {}'.format(epoch+1, num_epochs))
        print('-' * 10)

        for phase in ['train', 'validation']:
            if phase == 'train':
                model.train()
            else:
                model.eval()

            running_loss = 0.0
            running_corrects = 0

            for inputs, labels in dataloaders[phase]:
                inputs = inputs.to(device)
                labels = labels.to(device)

                outputs = model(inputs)
                loss = criterion(outputs, labels)

                if phase == 'train':
                    optimizer.zero_grad()
                    loss.backward()
                    optimizer.step()

                _, preds = torch.max(outputs, 1)
                running_loss += loss.item() * inputs.size(0)
                running_corrects += torch.sum(preds == labels.data)

            epoch_loss = running_loss / len(image_datasets[phase])
            epoch_acc = running_corrects.double() / len(image_datasets[phase])

            print('{0} loss: {1:.4f}, acc: {2:.4f}'.format(phase,
                                                          epoch_loss,
```

```
epoch_acc))  
return model
```

## ▼ Train model and save weights for future pre-trained usage

```
# Uncomment the below code lines in case if you want to train the RESNET50 model  
  
# model_trained = train_model(model, criterion, optimizer, num_epochs=3)  
  
# # # !mkdir models  
  
# torch.save(model_trained.state_dict(), '/content/drive/MyDrive/Colab Notebooks/models/weigh
```

## ▼ Load pre-train RESNET50

```
# Load model from weights when not training on the data.  
model.load_state_dict(torch.load('/content/drive/MyDrive/Colab Notebooks/models/weightsv1_2.h  
  
# Define Classifier  
classifier = NeuralNetClassifier(model, criterion=criterion)  
  
# Creating deep copies of original Classifier for Experimentation (Torch models/classifiers a  
classifier_random_sampling = copy.deepcopy(classifier)  
classifier_uncertainty = copy.deepcopy(classifier)  
classifier_entropy = copy.deepcopy(classifier)  
classifier_margin = copy.deepcopy(classifier)  
  
# Train loader to Array  
trainLoader = dataloaders['train']  
X, y = next(iter(trainLoader))  
X = X.detach().cpu().numpy()  
y = y.detach().cpu().numpy()  
  
print('Train data shape: ', X.shape, '\n',  
      'Number of classes: ', len(set(y)), '\n',  
      'all classes : ', trainLoader.dataset.dataset.class_to_idx)  
  
# Test loader to Array  
testLoader = dataloaders['test']  
X_test, y_test = next(iter(testLoader))  
X_test = X_test.detach().cpu().numpy()  
y_test = y_test.detach().cpu().numpy()  
  
print('\n\nTest data shape: ', X_test.shape, '\n',  
      'Number of classes: ', len(set(y)), '\n',
```

```
'all classes : ', testLoader.dataset.class_to_idx)

# Sampler loader to Array
samplerLoader = dataloaders['sampler']
X_sampler, y_sampler = next(iter(samplerLoader))
X_sampler = X_sampler.detach().cpu().numpy()
y_sampler = y_sampler.detach().cpu().numpy()

print('\n\nSampler data shape: ', X_sampler.shape, '\n',
      'Number of classes: ', len(set(y_sampler)), '\n',
      'all classes : ', samplerLoader.dataset.class_to_idx)

Train data shape: (250, 3, 128, 128)
Number of classes: 3
all classes : {'Bread': 0, 'Meat': 1, 'Vegetable-Fruit': 2}

Test data shape: (1031, 3, 128, 128)
Number of classes: 3
all classes : {'Bread': 0, 'Meat': 1, 'Vegetable-Fruit': 2}

Sampler data shape: (2250, 3, 128, 128)
Number of classes: 3
all classes : {'Bread': 0, 'Meat': 1, 'Vegetable-Fruit': 2}
```

## ▼ Define custom Query/Sampling functions for Active Learner (modAL)

```
# Function for random sampling (stratified) for Active learner query

def random_samp(clf, X_samp, y_samp, n_instances):
    # Randomly sample the remaining training data
    # INPUTS
    # clf : learner [modAL learner]
    # X_samp : [array] data to be queries or sampled
    # y_samp : [array] Targets of data to be sampled
    # n_instances : Number of images to be sampled
    # RETURNS [array] of indices of the queried X_samp, [array] Sampled output data from X_samp

    test_size_sampling = n_instances/len(X_samp) #Calculate the split ratio for the query
    _t, query_idx, _y, _ = train_test_split(range(len(X_samp)), y_samp, test_size=test_size_sampling)
    return query_idx, X_samp[query_idx]

# Function 2

# Function 3
```

## ▼ Define image visualization functions for use during Query experiment

```
# Display sample image predictions, n correct predictions and n wrong predictions

def display_predictions(nsamples, X_in, y_in, y_pred, classmap):
    correct = []
    wrong = []
    for i in range(len(y_in)):
        if y_in[i] == y_pred[i]:
            correct.append(i)
        else:
            wrong.append(i)

    # Choose random n samples for each correct/wrong predictions
    cor_idx = list(np.random.choice(correct, size=nsamples))
    wrn_idx = list(np.random.choice(wrong, size=nsamples))

    idx = cor_idx + wrn_idx
    arr = X_in.copy()
    # arr = np.reshape((3, 128,128), (128, 128, 3))
    arr = arr.swapaxes(1, 2)
    arr = arr.swapaxes(2, 3)
    fig, axes = plt.subplots(1, nsamples*2, figsize=(24,8))
    axes = axes.flatten()
    for i, ax in enumerate(axes):
        cnt = 1
        ind = idx[i]
        class_name = [name for name, num in classmap.items() if num == y_in[ind]][0]
        pred_name = [name for name, num in classmap.items() if num == y_pred[ind]][0]
        # print("\nGround Truth = ", class_name, " Prediction = ", pred_name)
        num_cols = 1
        num_rows = 1
        scale = 2
        Image = np.clip(np.array(((arr[ind]+1)/2)*255, np.int32), 0, 255)
        ax.imshow(Image)
        title = "GTruth:" + class_name + "\nPred:" + pred_name
        ax.title.set_text(title)
        # ax.titlesize(2)
        # ax = plt.subplot(nsamples*2, 1, cnt)
        # plt.imshow(arr[i], cmap=' ')
        cnt+=1
    return None

# Display top 5 and tail 5 images while looping Queries
warnings.filterwarnings("ignore", module = "matplotlib\..*")
```

```
def query_headtail(X_in, y_in, classmap, ntop, ntail):
    idx = list(range(ntop)) + list(range(-ntail, 0))
    arr = X_in.copy()
    # arr = np.reshape((3, 128,128), (128, 128, 3))
    arr = arr.swapaxes(1, 2)
```

```

arr = arr.swapaxes(2, 3)
fig, axes = plt.subplots(1, 10, figsize=(24,8))
axes = axes.flatten()
for i, ax in enumerate(axes):
    # for i in idx:
    cnt = 0
    ind = idx[i]
    class_name = [name for name, num in classmap.items() if num == y_in[ind]][0]
    # print("\nGround Truth = ", class_name)
    num_cols = 1
    num_rows = 1
    scale = 2
    # plt.subplot(2,10,cnt+1)    # the number of images in the grid is 5*5 (25)
    Image = np.clip(np.array(((arr[ind]+1)/2)*255, np.int32), 0, 255)
    ax.imshow(Image)
    ax.title.set_text(class_name)
    # ax.axis('off')
    # ax = plt.subplot(nsamples*2, 1, cnt)
    # plt.imshow(arr[i], cmap='')
    cnt+=1
return fig.show()

```

## ▼ Initialize modAL Learner for Base model

```

# initialize ActiveLearner
learner = ActiveLearner(
    estimator = classifier_random_sampling,
    X_training = X, y_training = y
    , query_strategy=random_samp,
)

class_map = trainLoader.dataset.dataset.class_to_idx

      epoch      train_loss      valid_acc      valid_loss      dur
      -----      -----      -----      -----      -----
      1          0.9375        0.5800        0.8269      15.9385
      2          0.8771        0.7600        0.7890      15.4435
      3          0.8390        0.9000        0.7700      15.4500
      4          0.8119        0.8800        0.7600      15.3978
      5          0.7903        0.8600        0.7521      15.4736
      6          0.7713        0.8000        0.7451      15.5910
      7          0.7531        0.7800        0.7384      15.3451
      8          0.7355        0.7800        0.7314      15.3260
      9          0.7181        0.8000        0.7235      15.2957
     10          0.7011        0.8400        0.7150      15.4077

```

## ▼ Initial Base Model Performance

```
# Predict on test split using initial model learner
```

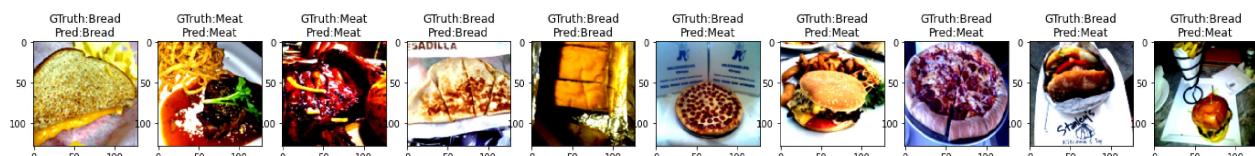
```
pred1 = learner.predict(X_test)
```

```
print(classification_report(y_test, pred1))
```

	precision	recall	f1-score	support
0	0.93	0.34	0.49	368
1	0.60	0.99	0.75	432
2	0.96	0.79	0.86	231
accuracy			0.71	1031
macro avg	0.83	0.70	0.70	1031
weighted avg	0.80	0.71	0.68	1031

```
# Display test data predictions
```

```
display_predictions(5, X_test, y_test, pred1, class_map)
```



## ▼ EXPERIMENT 1

### Random Sampling on Resnet50 pre-trained

```
# Active Learning on pre-trained Resnet50 with 250 new random sampled images from labeled dat
```

```
learner_randomsampling = copy.deepcopy(learner) # Deep copy of original learner
```

```
# Number of total instances to be sampled
```

```
num_instances = 250
```

```
# Random sampling Query
```

```
query_idx, query_instance = learner_randomsampling.query(X_sampler, y_sampler, num_instances)
print("ids: ",query_idx, "Instances: ", len(query_instance))
```

```
# Teach base model
```

```
learner_randomsampling.teach(
```

```
X=X_sampler[query_idx], y=y_sampler[query_idx], only_new=False,
)
```

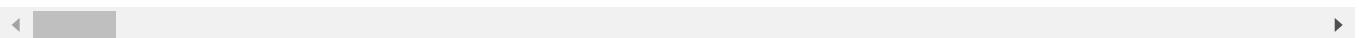
ids: [91, 2170, 372, 261, 888, 1679, 1044, 2095, 1994, 1428, 502, 1872, 2149, 148, 115]

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

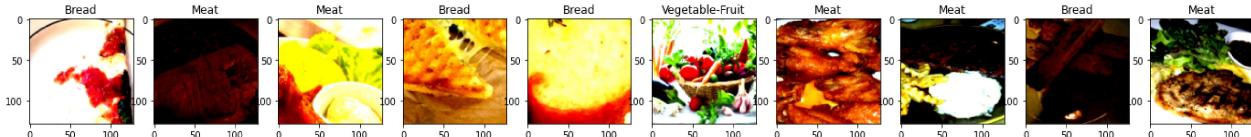
epoch	train_loss	valid_acc	valid_loss	dur
1	0.7266	0.7500	0.6927	30.3187
2	0.7062	0.7500	0.6771	30.5113
3	0.6835	0.7600	0.6608	30.3657
4	0.6611	0.7600	0.6452	30.7584
5	0.6394	0.7700	0.6294	30.5938
6	0.6186	0.7800	0.6143	30.5040
7	0.5988	0.7800	0.5998	30.6114
8	0.5798	0.7900	0.5861	30.7210
9	0.5618	0.7900	0.5734	30.7793
10	0.5448	0.8000	0.5611	30.9169



## ▼ Printing top 5 and last 5 images from the Queried data

```
print(query_headtail(X_sampler, y_sampler, class_map, 5, 5))
```

None



## ▼ Random Sampling Performance

```
pred0 = learner_randomsampling.predict(X_test)
print(classification_report(y_test, pred0))
```

	precision	recall	f1-score	support
0	0.95	0.39	0.56	368
1	0.61	0.99	0.75	432
2	0.98	0.74	0.85	231
accuracy			0.72	1031
macro avg	0.85	0.71	0.72	1031

weighted avg	0.82	0.72	0.70	1031
--------------	------	------	------	------

## ▼ EXPERIMENT 2

---

### Initialize modAL for Uncertainty Sampling

```
# initialize ActiveLearner 1 using pre-trained RESNET50 original classifier
learner1 = ActiveLearner(
    estimator = classifier_uncertainty,
    X_training = X, y_training = y
    , query_strategy=uncertainty_sampling
)
```

epoch	train_loss	valid_acc	valid_loss	dur
1	0.9375	0.5800	0.8269	15.1159
2	0.8771	0.7600	0.7890	15.1730
3	0.8390	0.9000	0.7700	15.2002
4	0.8119	0.8800	0.7600	15.1407
5	0.7903	0.8600	0.7521	15.2765
6	0.7713	0.8000	0.7451	15.2772
7	0.7531	0.7800	0.7384	15.0318
8	0.7355	0.7800	0.7314	15.1531
9	0.7181	0.8000	0.7235	15.1058
10	0.7011	0.8400	0.7150	15.1634

### ▼ Uncertainty Sampling on Resnet50 pre-trained

```
# Copy sampler data
X_sampler_copy = X_sampler.copy()
y_sampler_copy = y_sampler.copy()

# Number of Queries and instances
n_queries = 5
num_instances = 50

# Running Queries
for idx in range(n_queries):
    print('Query no. %d' % (idx + 1))
    # Uncertainty sampling query
    query_idx, query_instance = learner1.query(X_sampler_copy, n_instances=num_instances)
    print("ids: ",query_idx, "Instances: ", len(query_instance))
    print("Printing top and bottom Queried images: \n")
    print(query_headtail(X_sampler_copy[query_idx], y_sampler_copy[query_idx], class_map, 5,
```

```
# Teach base model
learner1.teach(
    X=X_sampler_copy[query_idx], y=y_sampler_copy[query_idx], only_new=False,
)
y_prediction = learner1.predict(X_test)
print("Report after ", num_instances, " new instances:\n", classification_report(y_test,
# remove queried instance from sampler pool
X_sampler_copy = np.delete(X_sampler_copy, query_idx, axis=0)
y_sampler_copy = np.delete(y_sampler_copy, query_idx, axis=0)
print(X_sampler_copy.shape)
```

Query no. 1

```
ids: [1570 1740 823 1054 2206 2202 1040 1422 2167 1634 1038 703 1630 911
180 1906 1066 1541 2101 971 1364 461 727 1955 2093 1502 680 1530
304 1394 755 134 873 1902 1648 35 295 1753 404 1459 1984 1086
175 671 2069 817 1368 1917 91 1501] Instances: 50
```

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.7751	0.8500	0.7139	18.2129
2	0.7568	0.8667	0.7033	18.0054
3	0.7445	0.8667	0.6932	18.1492
4	0.7326	0.8667	0.6831	18.1145
5	0.7207	0.8667	0.6731	17.8871
6	0.7090	0.8667	0.6635	17.7211
7	0.6974	0.8667	0.6541	17.5877
8	0.6860	0.8667	0.6447	17.7243
9	0.6748	0.8667	0.6354	17.6982
10	0.6638	0.8667	0.6262	17.7216

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.95	0.48	0.64	368
1	0.69	0.94	0.80	432
2	0.82	0.92	0.87	231
accuracy			0.77	1031
macro avg	0.82	0.78	0.77	1031
weighted avg	0.81	0.77	0.76	1031

(2200, 3, 128, 128)

Query no. 2

```
ids: [1335 1385 758 1866 295 78 660 305 1844 1835 1828 1883 68 324
639 66 253 1200 948 110 570 899 239 545 2039 1921 222 1929
1556 157 1356 721 1683 2156 1138 709 1566 1032 1045 855 1965 1332
1153 1680 1047 1380 1338 256 1716 1767] Instances: 50
```

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.7352	0.7857	0.6078	21.3312
2	0.7327	0.7571	0.6076	21.2800
3	0.7303	0.7714	0.5981	21.0662
4	0.7200	0.7714	0.5867	21.4384
5	0.7080	0.7714	0.5769	21.4085
6	0.6967	0.7857	0.5666	21.4945
7	0.6847	0.8000	0.5573	21.4028
8	0.6733	0.8000	0.5485	21.4899

9	0.6622	0.8000	0.5397	21.6291
10	0.6508	0.8143	0.5303	21.6469

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.68	0.96	0.80	368
1	0.95	0.63	0.76	432
2	0.91	0.87	0.89	231
accuracy			0.81	1031
macro avg	0.85	0.82	0.82	1031
weighted avg	0.85	0.81	0.80	1031

(2150, 3, 128, 128)

Query no. 3

ids: [1703 493 1082 1745 1519 1261 890 541 251 1404 1952 275 1718 294  
379 524 1176 1165 85 1485 827 331 1807 1904 1697 1701 1540 89  
1730 984 262 897 730 1468 447 1325 1262 539 394 1645 315 1480  
361 592 903 745 1208 2115 1251 274] Instances: 50

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.7042	0.8750	0.5047	24.2768
2	0.6597	0.8625	0.5049	24.3597
3	0.6496	0.8500	0.4978	24.3875
4	0.6366	0.8625	0.4900	24.1614
5	0.6236	0.8750	0.4825	24.1209
6	0.6106	0.8750	0.4755	24.3828
7	0.5979	0.8750	0.4691	24.2283
8	0.5856	0.8875	0.4630	24.2762
9	0.5735	0.8875	0.4572	24.1237
10	0.5617	0.8875	0.4516	24.1282

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.95	0.59	0.73	368
1	0.72	0.95	0.82	432
2	0.91	0.91	0.91	231
accuracy			0.81	1031
macro avg	0.86	0.82	0.82	1031
weighted avg	0.84	0.81	0.81	1031

(2100, 3, 128, 128)

Query no. 4

ids: [ 847 1883 398 1712 1140 1704 154 2005 823 905 672 74 1779 1531  
1472 1319 1309 1969 1014 1880 1012 44 1557 1299 29 1639 1396 464  
161 1967 1907 551 945 1002 626 1157 842 1863 1208 173 1943 1163  
996 277 437 1868 2038 1511 1677 617] Instances: 50

Printing top and bottom Queried images:

None

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.6228	0.8667	0.4205	28.0542
2	0.6063	0.8889	0.4133	27.9640
3	0.5964	0.8778	0.4068	28.0384
4	0.5867	0.8778	0.4009	27.8876
5	0.5772	0.8778	0.3954	27.8228
6	0.5679	0.8778	0.3902	27.8697
7	0.5587	0.8778	0.3852	27.7362
8	0.5497	0.8778	0.3804	28.0648
9	0.5409	0.8778	0.3759	27.7973
10	0.5323	0.8667	0.3716	27.8280

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.91	0.78	0.84	368
1	0.83	0.90	0.86	432
2	0.89	0.94	0.91	231
accuracy			0.87	1031
macro avg	0.87	0.87	0.87	1031
weighted avg	0.87	0.87	0.87	1031

(2050, 3, 128, 128)

Query no. 5

```
ids: [1402 1809 708 1798 1812 715 1516 2040 204 179 10 1013 1950 1245
    735 1870 311 1261 406 1912 688 1163 819 895 420 672 1308 1907
    1354 935 1655 135 261 1364 614 604 1693 1416 1692 101 1565 443
    770 153 527 1440 938 1390 33 578] Instances: 50
```

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

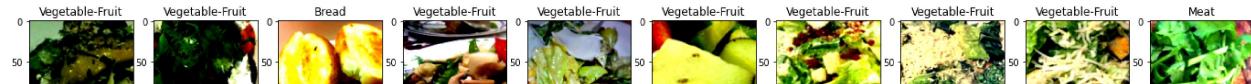
epoch	train_loss	valid_acc	valid_loss	dur
1	0.6084	0.8800	0.3710	29.7610
2	0.6013	0.8800	0.3668	29.8029
3	0.5897	0.8800	0.3636	29.8406
4	0.5787	0.8800	0.3607	29.8407
5	0.5684	0.8900	0.3576	29.7868
6	0.5585	0.8900	0.3549	29.8369
7	0.5491	0.8800	0.3523	29.6138
8	0.5400	0.8800	0.3493	29.3628
9	0.5312	0.8800	0.3465	29.6267
10	0.5227	0.8800	0.3438	29.5287

Report after 50 new instances:

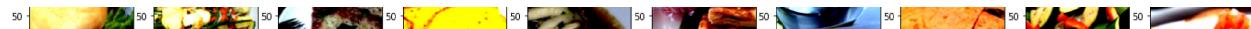
	precision	recall	f1-score	support
0	0.94	0.74	0.83	368
1	0.79	0.95	0.86	432
2	0.96	0.90	0.93	231

accuracy		0.87	1031
macro avg	0.90	0.87	1031
weighted avg	0.88	0.87	1031

(2000, 3, 128, 128)



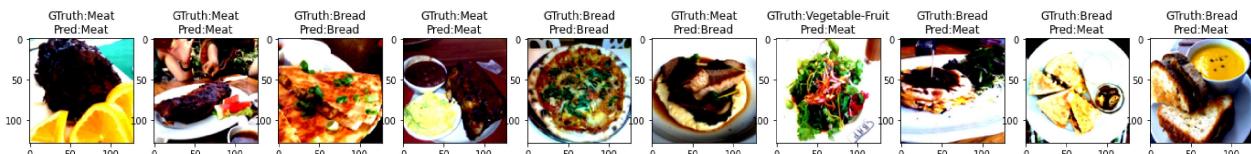
## ▼ Uncertainty Sampling Performance



```
pred2 = learner1.predict(X_test)
print(classification_report(y_test, pred2))
```

	precision	recall	f1-score	support
0	0.94	0.74	0.83	368
1	0.79	0.95	0.86	432
2	0.96	0.90	0.93	231
accuracy			0.87	1031
macro avg	0.90	0.87	0.87	1031
weighted avg	0.88	0.87	0.87	1031

```
# Display test data predictions
display_predictions(5, X_test, y_test, pred2, class_map)
```



## ▼ EXPERIMENT 3

### Initialize modAL for Entropy Sampling

```
# Initialize ActiveLearner 2 using pre-trained RESNET50 original classifier
learner2 = ActiveLearner(
```

```

estimator = classifier_entropy,
X_training = X, y_training = y
, query_strategy=entropy_sampling
)

```

epoch	train_loss	valid_acc	valid_loss	dur
1	0.9375	0.5800	0.8269	14.5356
2	0.8771	0.7600	0.7890	14.9374
3	0.8390	0.9000	0.7700	14.7878
4	0.8119	0.8800	0.7600	14.7829
5	0.7903	0.8600	0.7521	14.6104
6	0.7713	0.8000	0.7451	14.7384
7	0.7531	0.7800	0.7384	14.6705
8	0.7355	0.7800	0.7314	14.8638
9	0.7181	0.8000	0.7235	14.6466
10	0.7011	0.8400	0.7150	14.7210

## ▼ Entropy Sampling on Resnet50 pre-trained

```

# Copy sampler data
X_sampler_copy = X_sampler.copy()
y_sampler_copy = y_sampler.copy()

# Number of queries and instances
n_queries = 5
num_instances = 50

# Running Queries
for idx in range(n_queries):
    print('Query no. %d' % (idx + 1))
    # Entropy sampling query
    query_idx, query_instance = learner2.query(X_sampler_copy, n_instances=num_instances)
    print("ids: ",query_idx, "Instances: ", len(query_instance))
    print("Printing top and bottom Queried images: \n")
    print(query_headtail(X_sampler_copy[query_idx], y_sampler_copy[query_idx], class_map, 5,
    # Teach base model
    learner2.teach(
        X=X_sampler_copy[query_idx], y=y_sampler_copy[query_idx], only_new=False,
    )
    y_prediction = learner2.predict(X_test)
    print("Report after ", num_instances, " new instances:\n", classification_report(y_test,
    # remove queried instance from pool
    X_sampler_copy = np.delete(X_sampler_copy, query_idx, axis=0)
    y_sampler_copy = np.delete(y_sampler_copy, query_idx, axis=0)
    print(X_sampler_copy.shape)

```

Query no. 1

```
ids: [1634 727 1502 1459 1917 2167 911 1906 1902 1530 1394 1955 755 1038
1823 1040 16 1541 2069 180 1054 703 876 971 1422 1570 2202 854
1753 1066 823 1740 2007 2206 2093 1630 1718 1368 304 680 1633 1364
298 295 2195 1524 404 886 945 1086] Instances: 50
```

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.7769	0.8500	0.7144	17.8004
2	0.7593	0.8667	0.7040	17.9130
3	0.7472	0.8667	0.6941	17.7370
4	0.7356	0.8667	0.6845	17.7336
5	0.7240	0.8667	0.6749	17.8402
6	0.7125	0.8667	0.6652	18.1309
7	0.7011	0.8667	0.6557	17.8113
8	0.6898	0.8667	0.6463	17.7133
9	0.6787	0.8667	0.6371	17.9207
10	0.6678	0.8667	0.6280	17.8575

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.94	0.49	0.64	368
1	0.69	0.94	0.80	432
2	0.82	0.91	0.86	231
accuracy			0.77	1031
macro avg	0.82	0.78	0.77	1031
weighted avg	0.81	0.77	0.76	1031

(2200, 3, 128, 128)

Query no. 2

```
ids: [1828 1358 1965 111 762 307 6 1929 1140 1558 78 329 297 733
1568 1588 773 902 1921 2156 1890 663 1883 256 66 1049 1684 548
859 242 2104 492 665 1155 34 713 465 1866 1382 225 1047 642
1844 170 174 1835 1334 739 163 1954] Instances: 50
```

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.7349	0.8714	0.6084	20.7132
2	0.7261	0.8571	0.6001	20.8109
3	0.7182	0.8571	0.5887	21.9703
4	0.7066	0.8857	0.5772	21.0422
5	0.6942	0.9000	0.5664	20.9914
6	0.6820	0.9143	0.5563	20.9370
7	0.6700	0.9143	0.5461	20.9144
8	0.6577	0.9000	0.5368	20.9972

9	0.6462	0.9143	<b>0.5274</b>	20.9920
10	0.6345	0.9143	<b>0.5187</b>	20.8456

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.74	0.91	0.82	368
1	0.90	0.72	0.80	432
2	0.90	0.90	0.90	231
accuracy			0.83	1031
macro avg	0.85	0.85	0.84	1031
weighted avg	0.84	0.83	0.83	1031

(2150, 3, 128, 128)

Query no. 3

ids: [ 894 880 1747 379 1723 2018 294 1597 1261 1373 540 252 538 1470  
 226 766 1720 1487 1934 705 93 427 276 1319 1080 1705 154 1646  
 1521 394 860 141 84 331 729 1952 1406 1906 744 1884 1924 1175  
 621 1469 1124 1482 1490 696 1803 1326] Instances: 50

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.6935	<b>0.8625</b>	<b>0.5058</b>	23.5196
2	0.6660	0.8250	<b>0.5045</b>	23.8849
3	0.6567	0.8500	<b>0.4967</b>	23.7237
4	0.6449	0.8500	<b>0.4883</b>	23.8123
5	0.6329	0.8500	<b>0.4803</b>	23.7810
6	0.6211	0.8500	<b>0.4725</b>	23.9005
7	0.6094	0.8500	<b>0.4653</b>	23.6526
8	0.5979	0.8500	<b>0.4583</b>	23.4743
9	0.5867	0.8375	<b>0.4517</b>	23.7444
10	0.5757	0.8375	<b>0.4454</b>	23.6516

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.96	0.60	0.74	368
1	0.73	0.95	0.83	432
2	0.90	0.93	0.91	231
accuracy			0.82	1031
macro avg	0.86	0.83	0.82	1031
weighted avg	0.85	0.82	0.81	1031

(2100, 3, 128, 128)

Query no. 4

ids: [1281 903 881 1907 1475 952 596 844 1881 996 779 399 12 1036  
 1717 1399 414 420 1642 1140 675 1324 73 1164 1304 1190 42 1737  
 1210 280 156 832 1688 1533 1330 1370 1368 620 1784 27 1884 1002  
 9 24 109 43 542 305 1278 1698] Instances: 50

Printing top and bottom Queried images:

None

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.6232	0.8778	0.4178	26.7180
2	0.6032	0.8778	0.4115	27.0241
3	0.5928	0.8778	0.4058	26.9208
4	0.5826	0.8778	0.4003	26.8576
5	0.5725	0.8778	0.3951	26.8933
6	0.5627	0.8778	0.3901	26.8066
7	0.5530	0.8778	0.3852	26.8848
8	0.5435	0.8778	0.3805	26.9293
9	0.5341	0.8778	0.3760	26.8437
10	0.5250	0.8778	0.3717	26.8632

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.91	0.78	0.84	368
1	0.82	0.92	0.87	432
2	0.92	0.93	0.92	231
accuracy			0.87	1031
macro avg	0.88	0.88	0.88	1031
weighted avg	0.88	0.87	0.87	1031

(2050, 3, 128, 128)

Query no. 5

ids: [ 613 1869 1362 915 1473 1829 1496 2047 920 402 934 1260 407 414  
 1806 422 1512 161 257 1627 1988 1351 684 2040 951 480 174 1290  
 28 1760 785 1920 1865 1161 1647 1562 776 1702 1904 293 1130 1128  
 1599 1949 1025 731 578 593 986 71] Instances: 50

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

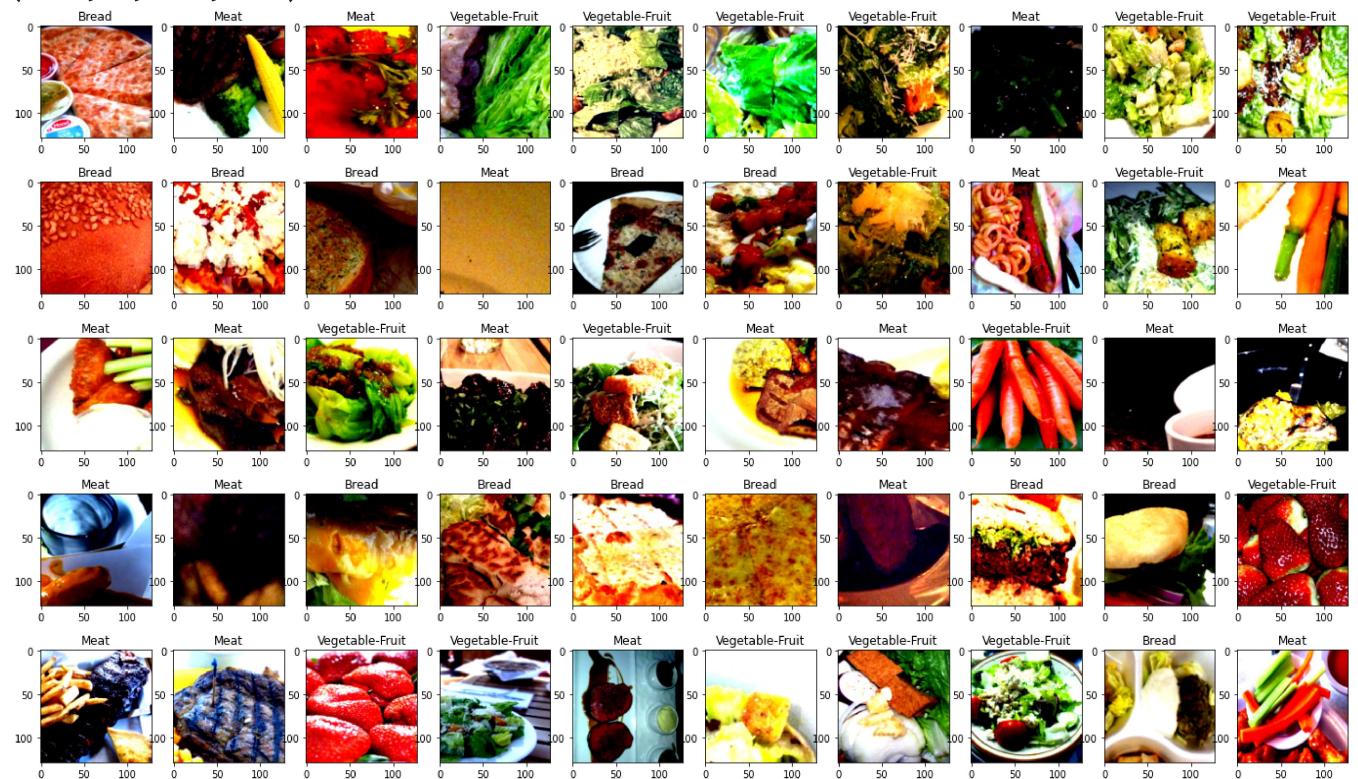
epoch	train_loss	valid_acc	valid_loss	dur
1	0.6008	0.8800	0.4181	29.4204
2	0.6250	0.8800	0.4143	29.3450
3	0.6135	0.8800	0.4101	29.4477
4	0.6019	0.8800	0.4068	29.4248
5	0.5908	0.8800	0.4027	29.3034
6	0.5796	0.8800	0.3990	29.2125
7	0.5689	0.8800	0.3949	29.3826
8	0.5582	0.8600	0.3901	29.2755
9	0.5474	0.8700	0.3862	29.5191
10	0.5373	0.8700	0.3816	29.1666

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.94	0.71	0.81	368
1	0.84	0.87	0.85	432
2	0.73	0.97	0.83	231

accuracy		<b>0.83</b>	<b>1031</b>
macro avg	<b>0.84</b>	<b>0.83</b>	<b>1031</b>
weighted avg	<b>0.85</b>	<b>0.83</b>	<b>1031</b>

(2000, 3, 128, 128)

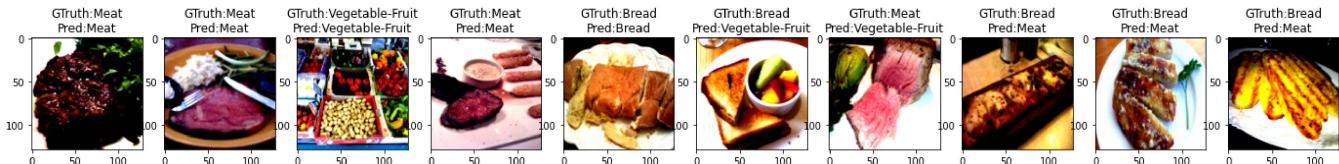


## ▼ Entropy Sampling Performance

```
y_pred2_2 = learner2.predict(X_test)
print(classification_report(y_test, y_pred2_2))
```

	precision	recall	f1-score	support
0	0.94	0.71	0.81	368
1	0.84	0.87	0.85	432
2	0.73	0.97	0.83	231
accuracy			0.83	1031
macro avg	0.84	0.85	0.83	1031
weighted avg	0.85	0.83	0.83	1031

```
display_predictions(5, X_test, y_test, y_pred2_2, class_map)
```



## ▼ EXPERIMENT 4

### Initialize modAL for Margin Sampling

```
# initialize ActiveLearner 3 using pre-trained RESNET50 original classifier
learner3 = ActiveLearner(
    estimator = classifier_margin,
    X_training = X, y_training = y
    , query_strategy=margin_sampling
)
```

epoch	train_loss	valid_acc	valid_loss	dur
1	0.9375	0.5800	0.8269	14.9222
2	0.8771	0.7600	0.7890	14.7982
3	0.8390	0.9000	0.7700	14.8877
4	0.8119	0.8800	0.7600	14.9304
5	0.7903	0.8600	0.7521	14.9605
6	0.7713	0.8000	0.7451	14.9707
7	0.7531	0.7800	0.7384	14.9900
8	0.7355	0.7800	0.7314	14.9430
9	0.7181	0.8000	0.7235	14.9480
10	0.7011	0.8400	0.7150	14.8245

## ▼ Margin Sampling on Resnet50 pre-trained

```
# Copy sampler data
X_sampler_copy = X_sampler.copy()
y_sampler_copy = y_sampler.copy()

# Number of Queries and instances
n_queries = 5
num_instances = 50

# Running Queries
for idx in range(n_queries):
    print('Query no. %d' % (idx + 1))
    # Margin sampling Query
    query_idx, query_instance = learner3.query(X_sampler_copy, n_instances = num_instances)
    print("ids: ",query_idx, "Instances: ", len(query_instance))
    print("Printing top and bottom Queried images: \n")
    print(query_headtail(X_sampler_copy[query_idx], y_sampler_copy[query_idx], class_map, 5,
    # Teach base model
    learner3.teach(
        X=X_sampler_copy[query_idx], y=y_sampler_copy[query_idx], only_new=False,
    )
    y_prediction = learner3.predict(X_test)
    print("Report after ", num_instances, " new instances:\n", classification_report(y_test,
    # remove queried instance from pool
    X_sampler_copy = np.delete(X_sampler_copy, query_idx, axis=0)
    y_sampler_copy = np.delete(y_sampler_copy, query_idx, axis=0)
    print(X_sampler_copy.shape)
```

Query no. 1

```
ids: [ 700 1878 1929 1038 200 1791 1597 1835 2220 484 1648 1194 2038 388
      369 1461 373 1333 2019 296 1972 1040 616 1321 643 1553 821 384
      1167 1515 285 1387 461 967 12 1643 1634 911 2101 428 763 180
      1688 554 1898 1564 1910 1984 1068 118] Instances: 50
```

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.7374	0.8667	0.7349	17.5197
2	0.7162	0.8667	0.7221	17.7790
3	0.7013	0.8833	0.7112	17.7133
4	0.6878	0.8833	0.7008	17.6944
5	0.6746	0.8833	0.6907	17.5841
6	0.6616	0.8833	0.6808	17.7433
7	0.6489	0.8833	0.6710	17.6901
8	0.6363	0.8833	0.6615	17.4559
9	0.6239	0.8833	0.6521	17.6547
10	0.6118	0.8833	0.6428	17.7647

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.84	0.85	0.84	368
1	0.86	0.87	0.86	432
2	0.91	0.88	0.89	231
accuracy			0.86	1031
macro avg	0.87	0.87	0.87	1031
weighted avg	0.86	0.86	0.86	1031

(2200, 3, 128, 128)

Query no. 2

```
ids: [ 569 1441 1417 1043 1061 2197 192 1874 1599 2021 719 1681 1864 41
      2166 636 608 576 114 886 1329 1772 1367 976 2190 783 1540 311
      1584 745 881 734 1379 916 611 623 631 2058 902 236 832 487
      1979 178 1025 248 2146 1647 710 2060] Instances: 50
```

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.6637	0.6571	0.6622	20.5377
2	0.6821	0.6286	0.6759	20.5511
3	0.6856	0.6286	0.6687	20.5884
4	0.6761	0.6286	0.6589	20.4439
5	0.6645	0.6571	0.6494	20.6948
6	0.6530	0.6714	0.6405	20.6633
7	0.6417	0.6714	0.6323	20.3725
8	0.6307	0.6857	0.6245	20.5627

9	0.6199	0.6857	<b>0.6173</b>	20.6097
10	0.6095	0.7143	<b>0.6104</b>	20.5680

Report after 50 new instances:

	precision	recall	f1-score	support
0	1.00	0.24	0.38	368
1	0.55	1.00	0.71	432
2	0.99	0.70	0.82	231
accuracy			0.66	1031
macro avg	0.85	0.64	0.64	1031
weighted avg	0.81	0.66	0.62	1031

(2150, 3, 128, 128)

Query no. 3

ids: [ 915 69 1666 1268 1897 671 1889 1696 1526 413 718 1966 720 324  
 721 1882 742 1386 753 1036 1017 1465 814 1853 592 1005 846 146  
 923 1816 2097 1226 1138 649 1451 1843 715 1800 941 1433 265 1224  
 1625 1375 1579 1991 1825 843 377 1560] Instances: 50

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.6054	<b>0.8500</b>	<b>0.5303</b>	23.6370
2	0.5552	0.8375	<b>0.5273</b>	23.7972
3	0.5501	0.8500	<b>0.5163</b>	23.8538
4	0.5391	<b>0.8625</b>	<b>0.5054</b>	23.7607
5	0.5278	0.8625	<b>0.4949</b>	23.4926
6	0.5164	<b>0.8750</b>	<b>0.4857</b>	23.8866
7	0.5056	0.8750	<b>0.4770</b>	23.6132
8	0.4950	<b>0.8875</b>	<b>0.4689</b>	23.6835
9	0.4847	0.8750	<b>0.4611</b>	23.8080
10	0.4746	0.8750	<b>0.4538</b>	23.7852

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.80	0.94	0.86	368
1	0.92	0.80	0.86	432
2	0.95	0.90	0.92	231
accuracy			0.87	1031
macro avg	0.89	0.88	0.88	1031
weighted avg	0.88	0.87	0.87	1031

(2100, 3, 128, 128)

Query no. 4

ids: [1392 1841 1720 392 1872 1691 747 298 1238 1921 1289 1588 159 361  
 893 457 665 911 1325 1552 1524 764 414 1655 1744 2 1611 1353  
 739 535 459 1577 550 1907 306 1919 1966 1766 95 2059 1264 86  
 612 1015 426 1539 1509 1881 630 829] Instances: 50

Printing top and bottom Queried images:

None

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.5082	0.8778	0.4227	26.5618
2	0.4756	0.8778	0.4165	27.0309
3	0.4647	0.8778	0.4101	27.0845
4	0.4548	0.8778	0.4040	26.8320
5	0.4452	0.8778	0.3982	26.8383
6	0.4359	0.8778	0.3926	27.2113
7	0.4268	0.8778	0.3873	26.8073
8	0.4180	0.8889	0.3822	26.8725
9	0.4094	0.8889	0.3774	26.8150
10	0.4011	0.8889	0.3727	26.8645

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.90	0.82	0.86	368
1	0.84	0.92	0.88	432
2	0.94	0.90	0.92	231
accuracy			0.88	1031
macro avg	0.89	0.88	0.89	1031
weighted avg	0.88	0.88	0.88	1031

(2050, 3, 128, 128)

Query no. 5

ids: [1131 96 1489 109 996 1503 135 1424 547 1921 977 1910 517 735  
 738 79 743 755 1377 1105 1375 2031 1644 1208 773 1367 1653 1657  
 76 229 239 64 1358 33 1357 421 419 281 1798 1270 1137 316  
 1718 376 826 1297 658 50 1307 1913] Instances: 50

Printing top and bottom Queried images:

None

Re-initializing module.

Re-initializing criterion.

Re-initializing optimizer.

epoch	train_loss	valid_acc	valid_loss	dur
1	0.4597	0.8200	0.4370	30.1028
2	0.4697	0.8200	0.4336	29.9555
3	0.4598	0.8100	0.4294	30.1222
4	0.4498	0.8000	0.4253	30.2353
5	0.4401	0.8000	0.4216	30.2481
6	0.4310	0.8000	0.4168	30.1675
7	0.4220	0.8000	0.4129	30.0994
8	0.4135	0.8000	0.4092	30.2309
9	0.4054	0.8000	0.4057	30.1239
10	0.3975	0.8000	0.4016	30.1725

Report after 50 new instances:

	precision	recall	f1-score	support
0	0.97	0.61	0.75	368
1	0.70	0.99	0.82	432
2	0.98	0.82	0.89	231

accuracy		<b>0.81</b>	<b>1031</b>
macro avg	<b>0.89</b>	<b>0.80</b>	<b>0.82</b>
weighted avg	<b>0.86</b>	<b>0.81</b>	<b>0.81</b>

(2000, 3, 128, 128)



## ▼ Margin Sampling Performance

```
y_pred3_2 = learner3.predict(X_test)
print(classification_report(y_test, y_pred3_2))
```

	precision	recall	f1-score	support
0	<b>0.97</b>	<b>0.61</b>	<b>0.75</b>	368
1	<b>0.70</b>	<b>0.99</b>	<b>0.82</b>	432
2	<b>0.98</b>	<b>0.82</b>	<b>0.89</b>	231
accuracy			<b>0.81</b>	<b>1031</b>
macro avg	<b>0.89</b>	<b>0.80</b>	<b>0.82</b>	<b>1031</b>
weighted avg	<b>0.86</b>	<b>0.81</b>	<b>0.81</b>	<b>1031</b>

```
display_predictions(5, X_test, y_test, y_pred3_2, class_map)
```

